

Ram Vamshi Krishna Rudraram

Data Engineer

MO | 417-419-8066 | rrudraram4@gmail.com

[Linkedin](#) | [Github](#) | [Portfolio](#)

SUMMARY

Data Engineer with 2 years building production ETL/ELT pipelines on AWS across 5+ TB of regulated healthcare data, plus an MS in Computer Science. Currently builds and operates a job-market data platform tracking 444,000+ postings ingested from roughly 22,900 sources. Depth in Python, PySpark, SQL, Airflow, and Redshift, with hands-on distributed processing, data quality, and pipeline orchestration.

TECHNICAL SKILLS

Languages: Python, SQL, PL/SQL, TypeScript, JavaScript, Bash

Data Engineering: Apache Spark, PySpark, Spark SQL & Streaming, Apache Airflow, dbt, Kafka, ETL/ELT, Dimensional Modeling, Star Schema, SCD Type 1/2, Data Quality, Data Governance, Schema Normalization

Cloud: AWS (S3, EMR, Glue, Redshift, Lambda, Athena, Kinesis, Step Functions); Azure (Data Factory, Databricks, Synapse, ADLS Gen2); DigitalOcean

Data Stores: Amazon Redshift, PostgreSQL (pgvector), MySQL, SQL Server, SQLite, Delta Lake, Parquet, Avro

Engineering & AI: Docker, Docker Compose, CI/CD, Git, Linux, Caddy, REST APIs, FastAPI, SQLAlchemy, Celery, Redis, Streamlit, pytest, TDD, Agile/Scrum; LLM Integration (OpenAI, Anthropic), RAG, Vector Embeddings, Pandas, NumPy, Power BI, Tableau

PROJECTS

Job Control Center | Job-market data platform, deployed and running 24/7

2026

Python, FastAPI, SQLAlchemy, SQLite (WAL), Streamlit, Docker Compose, Caddy, DigitalOcean

- Built a continuously running ingestion platform that crawls roughly 22,900 company career pages through 29 source-specific modules (Greenhouse, Lever, Ashby, Workday, iCIMS, SmartRecruiters, BambooHR, Paylocity, UKG, Oracle HCM, and job boards), normalizing heterogeneous payloads into a single schema across 444,000+ tracked postings.
- Diagnosed and fixed a source-fidelity defect in which a vendor's last-modified timestamp was ingested as the publish date, causing a 778-day-old posting to surface as new; enforced true publish-date semantics at the source layer and added explicit handling for the 43% of records whose providers emit no publish date.
- Designed a tiered retention and staleness model (90-day retention for high-value employers versus 10-day for others, with expiry detected from last-seen crawl timestamps rather than posting age) after short uniform windows were found to be dropping entire viable employers.
- Implemented a profile-matching scorer ranking every posting on role fit and employer signals, with a calibrated score threshold that suppresses high-volume duplicate-listing spam, plus a re-evaluation job to rescore stored rows when filter logic changes.
- Deployed a three-service stack (API, crawler, dashboard) via Docker Compose on a 2 vCPU / 2 GB droplet behind Caddy with HTTPS, applying CPU and memory limits to isolate the crawler from the API and tuning SQLite WAL access patterns around its single-writer constraint.

JobJarvis | github.com/ramvamshi019/jobjarvis

2026

FastAPI, Next.js, TypeScript, PostgreSQL/pgvector, Celery, Redis, dbt, Airflow, Docker

- Built a full-stack LLM job-search platform: async FastAPI backend (SQLAlchemy 2.0, Alembic) with Celery workers over Postgres and Redis, pgvector semantic matching against candidate profiles, dbt models, Airflow DAGs, and a Chrome extension, deployed via Docker Compose behind Caddy.

E-Commerce Support Classification | github.com/ramvamshi019/Ecommerce-ml-project

2025

PySpark, Python, Machine Learning, Distributed Computing (graduate coursework)

- Built an ML pipeline classifying e-commerce support cases and benchmarked distributed PySpark against single-node execution, quantifying runtime and scaling behavior across data volumes.

EXPERIENCE

Johnson & Johnson

Feb 2022 - Dec 2023

Data Engineer

Hyderabad, India

- Processed 5+ TB of clinical-trial and patient healthcare data by building batch and streaming ETL/ELT pipelines in Python, PySpark, and SQL on AWS (EMR, Glue, S3, Redshift) with Kafka and Kinesis ingestion of HL7 and FHIR records, delivering analytics-ready datasets to research teams.
- Cut query execution time on multi-billion-row tables by 50% by redesigning Redshift physical layout (distribution keys, sort keys, column compression) across star-schema data marts serving clinical reporting, regulatory compliance, and 100+ analysts.
- Reduced pipeline job failures by 35% by developing and maintaining 25+ Apache Airflow DAGs with retry, alerting, and SLA monitoring, improving orchestration reliability and observability.
- Accelerated new pipeline development by 30% team-wide by designing reusable, modular Python libraries for ingestion, transformation, logging, and validation, enforced through TDD and peer code review.

EDUCATION

Missouri State University

-Jan 2024 - Dec 2025

Master of Science, Computer Science

Springfield, MO

Koneru Lakshmaiah Education Foundation

- May 2023

Bachelor of Technology, Computer Science

India